

1 CS342: OS Lab
2 Department of CSE,
3 IIT, Guwahati, Assam 781 039

Exercise 03

4
5 OS Lessons: Process, Kernel data-structures, User and Kernel Stacks
6 Rating: Very hard
7

8 In this exercise, we shift our attention to Project 2: User Programs as described in Chapter 3 of
9 PintDoc. This is also one of the harder exercises as we need to find answers for many questions.

10 A few changes that you need to make to run pintos code for this new project:

11 The primary directory for User Programs project is `pintos/src/userprog`. To move to this
12 directory, we need to run `make clean` in directory `pintos/src/threads`. And, also
13 make the following changes in the script files:

14 The `perl` script `pintos` (which was placed previously in directory `~/bin`) be changed as
15 follows:

16 Line no.24: Replace "`$HOME/pintos/src/threads/build/os.dsk`" by
17 "`$HOME/pintos/src/userprog/build/os.dsk`"
18 (This file was mentioned in the installation instructions for Pintos)

19 Also, please change the last line in `Make.vars` file in directory
20 `pintos/src/userprog/userprog` to:
21 `SIMULATOR = --bochs`

22 This is also a good place to caution you about the small size of space available for the Kernel
23 stack(s). In fact, the page hosting the kernel stack for a thread also hosts all data-structures
24 required for the thread management by the kernel. You are also reminded that the kernel and
25 user virtual address spaces are disjoint and there are significant but essential controls built into
26 the software to monitor access across the boundary between these two regions.

27 Task for the Exercise:

28 A curt description of the task for the exercise is to write code to meet the specifications set in
29 PintDoc Section 3.3.3 *Argument Passing* on page 29. Let us be more friendly and supportive. It
30 would be helpful if you also read the document [Executing main\(\) function on Linux](#).

The first dot-point listed in Section 3.2 (page 28) suggests that we must write code in function `setup_stack()`. Section 3.1.4.1 on page 26 in PintDoc describes how the initial stack for function `main()` of the user program is organised. You must also carefully read the example in Section 3.5.1 *Program Startup Details* on pages 36-37.

We have already seen a function to print the contents of this initial stack in a previous exercise. The function is again listed here:

```
void test_stack(int *t)
{
    int i;
    int argc = t[1];
    char ** argv;

    argv = (char **) t[2];
    printf("ARGC:%d  ARGV:%x\n", argc, (unsigned int)argv);
    for (i = 0; i < argc; i++)
        printf("Argv[%d] = %x pointing at %s\n",
               i, (unsigned int)argv[i], argv[i]);
}
```

Now the question we must ask and answer is when can we call this function to print the contents of the activation record that function `main()` will receive?

There are two ways to start a user program in PintOS. One is to write a command line and pass it to PintOS during the kernel load time. The kernel command directive `run` (see Section 3.1.2 *Using the File System*) takes a user command as its argument. The other way is through an already running user program; we will come back to this method in a later exercise.

A user command may have its arguments. This may make some commands multi-word commands. Quotes (") are used to group such commands as a single argument for directive `run`. PintOS kernel load time directives are similar to built-in commands in Unix/Linux shells/ See near the bottom of page 24 [PintDoc] to understand how a command is specified as a single argument to directive `run`. For this exercise, this is the only way we use to start a user program.

What resources a program needs to run? A program running on a computer is called a process. A process is made of many tangible and virtual components:

- A process needs a thread trace (= a linear sequence of memory addresses – some executed and others to be executed – listing every instruction run on the computer processor by the user program) to execute instructions,
- Memory pages to store program instructions,
- Initialized and uninitialized data segments in memory (also called static data area), and
- A user stack (see Section 3.1.4.1 on page 26).

A user program is ready to run when these resources are available to the process. As we know from a previous exercise, all ready to run processes are placed in `ready_list` to be scheduled by the PintOS scheduler. From the time the `run` directive notes a command to be run to the time it can assemble all resources needed by the process, the process is in the state preceding `THREAD_READY`. After PintOS kernel has assembled a process with all the needed resources, the newly created process is inserted into `ready_list`. This occurs near label `done` in function `load()` in file `pintos/src/userprog/process.c`. This obviously is the place to ensure that we have set the program stack correctly. Consider if you wish to call function `test_stack()` near label `done`.

Let us now consider the activity history of a thread as the control arrives at label `done`. The directive `run` noted the user command to load and run in function `run_actions()` of file `pintos/src/threads/init.c`. We need to get a good understanding of the path from `run_action()` to `done` as the user command arguments must flow along this path to finally land in the run-time user stack to be available to function `main()` of the user program.

Use `cscope`, `ctags`, `gdb` and actual reading of the code to understand how the command made of a command-file name and its arguments is passed (or could be passed) from a calling function in the PintOS kernel code to the called function. Your path begins at function `run_action()` and ends at function `setup_stack()`. You may like to read *Exercise 06* document to get a better view of the mysterious tricks used to let the user program load itself into the memory.

Once you have understood the route to carry the required information from the command to function `setup_stack()` in file `process.c`, you are ready to write code in function `setup_stack()` to build the stack that will be made available to the user program when the user program begins execution. The new user program begins at the start of its function `main()`.

However, testing this user program is not possible at this stage! The reason for this limitation is lack of system calls to write messages on the computer console! Only kernel code can print messages; user programs cannot write on console yet.

97 We overcome this limitation by calling function `test_stack()` in the kernel code (and not in
98 the user program code).

99 How we tested our implementation:

100 The test sequence below is a reconstruction of the original activities that we used when we first
101 completed this exercise. (The Pintos code being used to re-enact the exercise for this section is
102 the completed code for the project. This completed code performance many more tasks.)
103 Caution: You may experience significant differences in your output.

104 In the script below, the text typed by the user has been shown in bold. Some output from the
105 standard Pintos utilities has been deleted as it provides no useful insight. The output that is of
106 minor interest is shown in smaller fonts to fit the page width neatly.

107

108 These commands were used to compile programs in directory `~/pintos/src/examples`

109 `[vmm@progsrv ~]$ cd pintos/src/examples/`

110 `[vmm@progsrv examples]$ make`

111 [Output deleted]

112 These commands were used to setup Pintos to load and run a user program.

113 `[vmm@progsrv examples]$ cd ../userprog/`

114 `[vmm@progsrv userprog]$ make`

115 `cd build && make all`

116 `make[1]: Entering directory`

117 ``/home/CS342/2016/FAC/vmm/pintos/src/userprog/build'`

118 `make[1]: Nothing to be done for `all'.`

119 `make[1]: Leaving directory`

120 ``/home/CS342/2016/FAC/vmm/pintos/src/userprog/build'`

121 `[vmm@progsrv userprog]$ cd build/`

122 `[vmm@progsrv build]$ pintos-mkdisk fs.dsk 2`

123 `[vmm@progsrv build]$ pintos -q -f`

124 [Output deleted]

125 Finally, we test our implementation of function `stack_setup()` :

126 `[vmm@progsrv build]$ pintos -p ../../examples/echo -a echo -- -q`

127 `[vmm@progsrv build]$ pintos -q run "echo My stack_setup() works"`

```

128 Writing command line to /tmp/EHglakwWBy.dsk...
129 squish-ptty bochs -q
130 =====
131                Bochs x86 Emulator 2.5.1
132                Built from SVN snapshot on January 6, 2012
133                Compiled on Oct 10 2012 at 11:12:02
134 =====
135 000000000000i[      ] reading configuration from bochsrc.txt
136 000000000000i[      ] installing nogui module as the Bochs GUI
137 000000000000i[      ] using log file bochsout.txt
138 Kernel command line: -q run 'echo My stack_setup() works'
139 Pintos booting with 4,096 kB RAM...
140 370 pages available in kernel pool.
141 369 pages available in user pool.
142 Calibrating timer... 204,600 loops/s.
143 hd0:0: detected 1,008 sector (504 kB) disk, model "Generic 1234", serial
144 "BXHD00011"
145 hd0:1: detected 4,032 sector (1 MB) disk, model "Generic 1234", serial
146 "BXHD00012"
147 Boot complete.
148 Executing 'echo My stack_setup() works':
149 ARGV:4  ARGV:bfffffff04
150 Argv[0] = bfffffff0 pointing at echo
151 Argv[1] = bffffffed pointing at My
152 Argv[2] = bffffffdf pointing at stack_setup()
153 Argv[3] = bffffffd9 pointing at works
154
155 The last few lines above are of primary interest to verify the completion of the exercise. The
156 remaining output below is from Pintos code that has not yet been included in your project. You
157 may notice significant variation in your output (but the variation is not relevant to your
158 exercise.)
159
160 echo My stack_setup() works
161 echo: exit(0)
162 Execution of 'echo My stack_setup() works' complete.
163 Timer: 183 ticks

```

164 Thread: 0 idle ticks, 133 kernel ticks, 53 user ticks

165 hd0:0: 0 reads, 0 writes

166 hd0:1: 28 reads, 0 writes

167 Console: 819 characters output

168 Keyboard: 0 keys pressed

169 Exception: 0 page faults

170 Powering off...

171 =====

172 Bochs is exiting with the following message:

173 [UNMP] Shutdown port: shutdown requested

174 =====

175 [vmm@progsrv build]\$

176

177 **Contributing Authors:**

178 Vishv Malhotra, Gautam Barua, Rashmi Dutta Baruah