

CS342: OS Lab

Department of CSE,

IIT, Guwahati, Assam 781 039

Exercise 04

OS Lessons: Interrupts, System calls, Exit Status, Return Value

Rating: Moderate

Please do not start on this exercise until you have successfully completed Exercise 03.

Attempting to progress without a good understanding of the solutions needed in the exercise is likely to be expensive on your time and efforts.

Task for the Exercise:

System calls are user program's requests to the kernel to perform some service. User programs are usually prevented from performing these activities directly because their actions may affect the other processes (user programs and even the kernel) running on the system. The kernel provides the coordinating authority to oversee these activities.

A consequence of this constrain is that the kernel must work in a robust isolation from the user program that has demanded the service. Thus, kernel and user stacks are separate entities.

At the same time, the kernel needs to receive system-call requests and their arguments to provide the requested services. Further, the kernel needs a way to communicate the results and service status information back to the user programs (or processes). So we need communication channels between the kernel and the user virtual address spaces. To better understand the issues at hand, these sections in PintDoc are definitely your essential readings:

- Section 3.1.5 *Accessing User Memory*,
- Section 3.4.2 *System Calls FAQ*,
- Section 3.5.2 *System Call Details*,
- Section 3.3.2 *Process Termination Messages*,
- Appendix A.4 *Interrupt Handling*,
- Appendix A.4.1 *Interrupt Infrastructure*,
- Appendix A4.2 *Internal Interrupt Handling*

It is important that you take advantage of the program pattern that PintOS code (See function `run_actions()` in file `threads/init.c`) uses to call functions. The pattern maintains an array of pointers to functions. The system call number is used as index in the array to jump to the appropriate function implementing the system call. This is not just efficient, it is less

vulnerable to errors and eases the task of adding new functions later. The pattern also keeps the implementation of each system call in a separate function. Function `userprog/syscall.c` referred to in Section 3.3.4 *System Calls* becomes a convenient arrival and departure point for all system calls.

Your first and very easy task is to create the skeleton functions for each of the system call mentioned in User Program project and provide a way to call the right function for each system call number for which a request arrives at function `syscall_handler()` in file `userprog/syscall.c`. Each of these functions can initially be set to mirror the older behavior of function `syscall_handler()`.

Goal for the exercise is to implement and have (at least) the following system calls (see section 3.3.4 *System Calls*) running: `halt()`, `exit()`, and `write()` on `stdout` (`fd = 1`). This will give you at least 10 working test cases on `make check` command.

The `write` system call requires your implementation to return a value back to the user program. The task needs an arrangement that would need a bit of your attention.

System call `halt()` is easy to implement once you work out a small correction you need to make in the specification given on page 29.

System call `exit()` requires whole of the advice given in Section 3.2 *Suggested Order of Implementation* implemented and a bit more.

Our advice is to implement function `process_wait()` in file `process.c` to work as follows:

While the thread associated with `child_tid` is not in state `THREAD_DYING` call `thread_yield()` to avoid spin wait. However, you need a small but subtle change in the provided kernel code to have this part working. This change relates to an allocated resource and when it is safe to release the resource. For the present we may decide to keep it instead of deallocating the resource.

In the last exercise of *User Program* project we will revisit the issue of the resource deallocation more earnestly. Unallocated (also called, leaked) resources degrade the Kernel's ability to host multiple threads/processes over time. Every allocated resource need to be deallocated if the kernel need to sustain vitality indefinitely.

Many a times, PintDoc makes an important suggestion without stressing its importance. The second paragraph after description of system call `close` on page 32 is a good advice. It says that the addresses provided to the system calls by the user program must be carefully checked. PintDoc recommendation is to do so now. The sanitation has two parts. First, determine that

67 the address is a valid user virtual address. Second part is to ensure that the address does
68 translate into a real physical address. If either condition is violated, the address is a likely source
69 of trouble for the kernel integrity and stability.

70 You must create useful functions in file `syscall.c` to test the validity of the addresses passed
71 as arguments in the system calls. The arguments to be tested for valid address and nature of
72 the test differs for each system call. It suffices at this stage to only exercise this test for the
73 system calls you have implemented in this exercise.

74 You may continue to implement other system calls related to files if you wish. Or you may wait
75 for the next exercise. We recommend that you leave system calls `exec()`, `wait()` for the last
76 exercise. This also may apply to the task described in Section 3.3.5 *Denying Writes to*
77 *Executables*.

78 Our experiences with `make check`:

79 Once again, we have tried to revert our implementation to the stage at the completion of this
80 exercise to give you an indication of what your program should be able to perform at this stage.

81 You must expect some small variations between your achievements and the one listed here.
82 However, we are in less than 50 failed tests range.

83

```
84 [vmm@progsrv build]$ make check
85 pass tests/userprog/args-none
86 pass tests/userprog/args-single
87 pass tests/userprog/args-multiple
88 pass tests/userprog/args-many
89 pass tests/userprog/args-dbl-space
90 pass tests/userprog/sc-bad-sp
91 pass tests/userprog/sc-bad-arg
92 pass tests/userprog/sc-boundary
93 pass tests/userprog/sc-boundary-2
94 pass tests/userprog/halt
95 pass tests/userprog/exit
96 FAIL tests/userprog/create-normal
97 pass tests/userprog/create-empty
98 FAIL tests/userprog/create-null
99 FAIL tests/userprog/create-bad-ptr
100 pass tests/userprog/create-long
101 FAIL tests/userprog/create-exists
102 FAIL tests/userprog/create-bound
103 FAIL tests/userprog/open-normal
104 FAIL tests/userprog/open-missing
105 FAIL tests/userprog/open-boundary
```

106 FAIL tests/userprog/open-empty
107 pass tests/userprog/open-null
108 FAIL tests/userprog/open-bad-ptr
109 FAIL tests/userprog/open-twice
110 FAIL tests/userprog/close-normal
111 FAIL tests/userprog/close-twice
112 pass tests/userprog/close-stdin
113 pass tests/userprog/close-stdout
114 pass tests/userprog/close-bad-fd
115 FAIL tests/userprog/read-normal
116 FAIL tests/userprog/read-bad-ptr
117 FAIL tests/userprog/read-boundary
118 FAIL tests/userprog/read-zero
119 pass tests/userprog/read-stdout
120 pass tests/userprog/read-bad-fd
121 FAIL tests/userprog/write-normal
122 FAIL tests/userprog/write-bad-ptr
123 FAIL tests/userprog/write-boundary
124 FAIL tests/userprog/write-zero
125 pass tests/userprog/write-stdin
126 pass tests/userprog/write-bad-fd
127 FAIL tests/userprog/exec-once
128 FAIL tests/userprog/exec-arg
129 FAIL tests/userprog/exec-multiple
130 FAIL tests/userprog/exec-missing
131 pass tests/userprog/exec-bad-ptr
132 FAIL tests/userprog/wait-simple
133 FAIL tests/userprog/wait-twice
134 FAIL tests/userprog/wait-killed
135 pass tests/userprog/wait-bad-pid
136 FAIL tests/userprog/multi-recurse
137 FAIL tests/userprog/multi-child-fd
138 FAIL tests/userprog/rox-simple
139 FAIL tests/userprog/rox-child
140 FAIL tests/userprog/rox-multichild
141 pass tests/userprog/bad-read
142 pass tests/userprog/bad-write
143 pass tests/userprog/bad-read2
144 pass tests/userprog/bad-write2
145 pass tests/userprog/bad-jump
146 pass tests/userprog/bad-jump2
147 FAIL tests/userprog/no-vm/multi-oom
148 FAIL tests/filesys/base/lg-create
149 FAIL tests/filesys/base/lg-full
150 FAIL tests/filesys/base/lg-random
151 FAIL tests/filesys/base/lg-seq-block
152 FAIL tests/filesys/base/lg-seq-random

```
153 FAIL tests/filesys/base/sm-create
154 FAIL tests/filesys/base/sm-full
155 FAIL tests/filesys/base/sm-random
156 FAIL tests/filesys/base/sm-seq-block
157 FAIL tests/filesys/base/sm-seq-random
158 FAIL tests/filesys/base/syn-read
159 FAIL tests/filesys/base/syn-remove
160 FAIL tests/filesys/base/syn-write
161 47 of 76 tests failed.
162 make: *** [check] Error 1
163 [vmm@progsrv build]$
164
```

165 **Contributing Authors:**

166 Vishv Malhotra, Gautam Barua, Rashmi Dutta Baruah