

4
5 OS Lessons: Exec(), Wait() and Denying Writes to Executables
6 Rating: Hard
7

8 Please do not start on this exercise before you have successfully completed Exercise 04.
9 Attempting to progress without a good understanding of the solutions needed in Exercise 04 is
10 likely to be expensive on your time and efforts. It is also a good idea to complete Exercise 05
11 first.

12 Tasks for the Exercise:

13 In this exercise, we work on system calls `exec()` and `wait()`. System call `exec()`
14 specifications may be modelled on function `process_execute()` that was written in an
15 earlier exercise. However, you must also make sure that your implementation does not report a
16 successful completion of the system call until after the child process/thread is a certain goer.

17 The latter requirement is not trivial. A new process can fail to reach a successful birth for many
18 reasons. Thus, a successful birth should be report at the completion of all stages; and, not at
19 the start when memory for `struct thread` is allocated.

20 PintDoc suggests that `wait()` implementation is a difficult task. The remark about the
21 difficulty is not without its merit.

22 Described below is a possible design suggestion that you may use for these two system calls. It
23 is not a praise-worthy plan but it worked for us. We used a set of 4 arrays indexed by thread
24 identifier. For each thread, the array-set records data that is not conveniently recorded in
25 `struct thread`. Example of information that is inconvenient to record in `struct`
26 `thread` is information that is needed even if the thread no longer exists or was not created
27 successfully.

28 Some details of our implementation of these 4 arrays is given below but we do expect that keen
29 students will improve on these and construct better data-structures to record data about each
30 thread that PintOS kernel needs outside the life-span of the threads. `Struct thread` is a
31 convenient data-structure to hold data needed during the active life-span of a thread.

- 32 1. One of the arrays we used provides mapping from thread-identifier (`tid`) to the
33 matching thread's `struct thread`.

2. The second array was used to record if the thread identified by `tid` has completed its birth successfully. Successful birth is achieved after the process's program code has been loaded. A `tid` is allocated as soon as space is allocated for data-structure `struct thread`. `PintDoc` requires that `exec()` does not return thread's `tid` before the thread is properly established.
3. The third array is used to park `exit-status` of a thread for system call `wait()` from the thread's parent.
4. Our final array helps in managing accesses to these data-structures during and after the thread's `THREAD_DYING` epoch.

Swiss Army Knife:

`Struct thread` is a veritable multi-purpose structure capable of supporting at least 3 different traces of program execution. In implementing `exec()`, the students need to understand the structure in its full details. Indeed, this sophistication is the reason, we delayed the implementation of `exec()` and `wait()` as the last item in this project. To help you understand the details, we list below function `thread_create()` from file `threads/thread.c`:

```
tid_t
thread_create (const char *name, int priority,
               thread_func *function, void *aux)
{
    struct thread *t;
    struct kernel_thread_frame *kf;
    struct switch_entry_frame *ef;
    struct switch_threads_frame *sf;
    tid_t tid;
    enum intr_level old_level;

    ASSERT (function != NULL);

    /* Allocate thread. */
    t = palloc_get_page (PAL_ZERO);
    if (t == NULL)
        return TID_ERROR;

    /* Initialize thread. */
    init_thread (t, name, priority);
    tid = t->tid = allocate_tid ();

    /* Prepare thread for first run by initializing its stack.
       Do this atomically so intermediate values for the 'stack'
       member cannot be observed. */
    old_level = intr_disable ();

    /* Stack frame for kernel_thread(). */
```

```

78     kf = alloc_frame (t, sizeof *kf);
79     kf->eip = NULL;
80     kf->function = function;
81     kf->aux = aux;
82
83     /* Stack frame for switch_entry(). */
84     ef = alloc_frame (t, sizeof *ef);
85     ef->eip = (void (*) (void)) kernel_thread;
86
87     /* Stack frame for switch_threads(). */
88     sf = alloc_frame (t, sizeof *sf);
89     sf->eip = switch_entry;
90     sf->ebp = 0;
91
92     intr_set_level (old_level);
93
94     /* Add to run queue. */
95     thread_unblock (t);
96
97     return tid;
98 }
99

```

100 A newly created thread, begins its life by executing function `kernel_thread()`. The
101 function is copied below to aid our explanation. Note in the code above, how a call to function
102 `kernel_thread()` is included in the code. This is not the way your first programming
103 course taught you! The function run will be initiated by the PintOS scheduler.

```

104 /* Function used as the basis for a kernel thread. */
105 static void
106 kernel_thread (thread_func *function, void *aux)
107 {
108     ASSERT (function != NULL);
109
110     intr_enable (); /* The scheduler runs with interrupts off. */
111     function (aux); /* Execute the thread function. */
112     thread_exit (); /* If function() returns, kill the thread. */
113 }
114

```

115 You need to recognize that the function finds its arguments on a stack created previously
116 during `thread_create()` execution. Argument `aux` is really the command line specifying
117 the user program with arguments to be run in the process being assembled. And, argument
118 `function`, is function `start_process()` – different from the user program you might
119 have expected!

120 Function `start_process()` is responsible to load the user program and setup user-
121 program's initial stack. All the needed information is in argument `aux`.

122 Once again we tell you that function `kernel_thread()` is run as a new entity (child thread)
123 separate from the thread that called function `thread_create()`. The separation occurred
124 near the end of the function `thread_create()` as should be noted through the code:

```
125     /* Add to run queue. */  
126     thread_unblock (t);  
127
```

128 For the sake of completeness we say, all context switch from a running thread to a new running
129 thread occur through function `switch_threads()`. Thus, the thread exiting status
130 `THREAD_RUNNING` and the thread entering status `THREAD_RUNNING` seamlessly execute the
131 same code in function `switch_threads()`. The first switch for a thread is special as it has
132 no past to resume from. Function `thread_create()` provides frame `switch_entry()`
133 and it is no surprise that the entry is the start of function `kernel_thread()`. This is the
134 function every thread runs at their birth.

135 It may be interesting and useful to read the appendix before reading this section again. But this
136 is a magic trick every good computer student must learn and master.

137 Now Enjoy Coding the Final Part of Project:

138 The specifications for system calls `exec()` and `wait()` are near replica of functions
139 `process_execute()` and `process_wait()` in file `userprog/process.c`. But, the
140 differences are important too.

141 We expect (and recommend) that you work on this implementation in three phases. In Phase 1,
142 focus on developing code to correctly implement system call `exec()`.

143 In the Phase2, include system call `wait()` into your goals.

144 In the final phase, you work on the specifications set in Section 3.3.5 *Denying Writes to*
145 *Executables* of PintDoc.

146 The success of each phase, as determined by the success status of command `make check`, is
147 given below to help you monitor your progress.

148 Status of our “make check”:

149

150 After full implementation of `exec()` :

151

```
152 [vmm@progsrv build]$ make check  
153 pass tests/userprog/args-none
```

154 pass tests/userprog/args-single
155 pass tests/userprog/args-multiple
156 pass tests/userprog/args-many
157 pass tests/userprog/args-dbl-space
158 pass tests/userprog/sc-bad-sp
159 pass tests/userprog/sc-bad-arg
160 pass tests/userprog/sc-boundary
161 pass tests/userprog/sc-boundary-2
162 pass tests/userprog/halt
163 pass tests/userprog/exit
164 pass tests/userprog/create-normal
165 pass tests/userprog/create-empty
166 pass tests/userprog/create-null
167 pass tests/userprog/create-bad-ptr
168 pass tests/userprog/create-long
169 pass tests/userprog/create-exists
170 pass tests/userprog/create-bound
171 pass tests/userprog/open-normal
172 pass tests/userprog/open-missing
173 pass tests/userprog/open-boundary
174 pass tests/userprog/open-empty
175 pass tests/userprog/open-null
176 pass tests/userprog/open-bad-ptr
177 pass tests/userprog/open-twice
178 pass tests/userprog/close-normal
179 pass tests/userprog/close-twice
180 pass tests/userprog/close-stdin
181 pass tests/userprog/close-stdout
182 pass tests/userprog/close-bad-fd
183 pass tests/userprog/read-normal
184 pass tests/userprog/read-bad-ptr
185 pass tests/userprog/read-boundary
186 pass tests/userprog/read-zero
187 pass tests/userprog/read-stdout
188 pass tests/userprog/read-bad-fd
189 pass tests/userprog/write-normal
190 pass tests/userprog/write-bad-ptr
191 pass tests/userprog/write-boundary
192 pass tests/userprog/write-zero
193 pass tests/userprog/write-stdin
194 pass tests/userprog/write-bad-fd
195 FAIL tests/userprog/exec-once
196 FAIL tests/userprog/exec-arg
197 FAIL tests/userprog/exec-multiple
198 pass tests/userprog/exec-missing
199 pass tests/userprog/exec-bad-ptr
200 FAIL tests/userprog/wait-simple
201 FAIL tests/userprog/wait-twice
202 FAIL tests/userprog/wait-killed
203 pass tests/userprog/wait-bad-pid
204 FAIL tests/userprog/multi-recurse
205 pass tests/userprog/multi-child-fd

```

206 FAIL tests/userprog/rox-simple
207 FAIL tests/userprog/rox-child
208 FAIL tests/userprog/rox-multichild
209 pass tests/userprog/bad-read
210 pass tests/userprog/bad-write
211 pass tests/userprog/bad-read2
212 pass tests/userprog/bad-write2
213 pass tests/userprog/bad-jump
214 pass tests/userprog/bad-jump2
215 FAIL tests/userprog/no-vm/multi-oom
216 pass tests/filesys/base/lg-create
217 pass tests/filesys/base/lg-full
218 pass tests/filesys/base/lg-random
219 pass tests/filesys/base/lg-seq-block
220 pass tests/filesys/base/lg-seq-random
221 pass tests/filesys/base/sm-create
222 pass tests/filesys/base/sm-full
223 pass tests/filesys/base/sm-random
224 pass tests/filesys/base/sm-seq-block
225 pass tests/filesys/base/sm-seq-random
226 FAIL tests/filesys/base/syn-read
227 pass tests/filesys/base/syn-remove
228 FAIL tests/filesys/base/syn-write
229 13 of 76 tests failed.
230 make: *** [check] Error 1
231
232
233
234 On completion of exec() and wait() system calls:

```

235 Our success status improved to just 4 fails:

```

236 pass tests/userprog/wait-twice
237 pass tests/userprog/wait-killed
238 pass tests/userprog/wait-bad-pid
239 pass tests/userprog/multi-recurse
240 pass tests/userprog/multi-child-fd
241 FAIL tests/userprog/rox-simple
242 FAIL tests/userprog/rox-child
243 FAIL tests/userprog/rox-multichild
244 pass tests/userprog/bad-read
245 pass tests/userprog/bad-write
246 pass tests/userprog/bad-read2
247 pass tests/userprog/bad-write2
248 pass tests/userprog/bad-jump
249 pass tests/userprog/bad-jump2
250 pass tests/userprog/no-vm/multi-oom
251 pass tests/filesys/base/lg-create
252 pass tests/filesys/base/lg-full
253 pass tests/filesys/base/lg-random
254 pass tests/filesys/base/lg-seq-block

```

```
255 pass tests/filesys/base/lg-seq-random
256 pass tests/filesys/base/sm-create
257 pass tests/filesys/base/sm-full
258 pass tests/filesys/base/sm-random
259 pass tests/filesys/base/sm-seq-block
260 pass tests/filesys/base/sm-seq-random
261 pass tests/filesys/base/syn-read
262 pass tests/filesys/base/syn-remove
263 FAIL tests/filesys/base/syn-write
264 4 of 76 tests failed.
265 make: *** [check] Error 1
```

266

267 On completion of all three phases:

268 And, finally after successful implementation of *Denying Writes to Executables* we could pass all
269 tests. This part does require a lot of thought and several score lines of kernel code.

270 The output below is also a confirmation that if a student meticulously follows the instructions,
271 all tests in *User Program* project can be successfully completed.

272 In nutshell, the implementation requires that for each executable file that has been loaded in
273 one or more active processes, we maintain a count of the processes that have loaded the same
274 executable file. The denial of write obligation on the executable file stays till the last process
275 loaded with the file has terminated. A smart student would have already sensed that it requires
276 very smart programming as we cannot know which process loaded with this executable file will
277 be the last to terminate. It also stands to reason that the process that loaded the executable file
278 first (and hence initiates the constraint “deny write on the file”), will be among the early one to
279 finish ahead of those who loaded the same file after it. The “deny write on file” constraint is
280 lifted as soon as this process terminates even though there may be other processes loaded with
281 the file still active in the system.

282 The problem described in the last paragraph is not the only tricky issue that you need to handle.
283 You also need to understand that not all threads are subject to a `wait()` call from their parent
284 thread. That is, `wait()` is not a reliable indicator of the termination of a process. However, all
285 resources given to a thread must be returned on its completion. That is, to avoid resource
286 leakage every page carrying a `struct thread` data-structure needs to be freed by calling
287 `palloc_free_page()`. Likewise, every open file must be closed. If you fail to deallocate all
288 resources, your kernel degrades over time. Your resource deallocation plan cannot rely on
289 function `wait()` to free resources.

290 The test `tests/userprog/no-vm/multi-oom` only succeeded when the
291 implementation supported 2040 threads each with ability to host 128 open files
292 simultaneously.

```
293 [vmm@progsrv ~]$ cd pintos/src/userprog/build/
294 [vmm@progsrv build]$ make check
295 pass tests/userprog/args-none
296 pass tests/userprog/args-single
297 pass tests/userprog/args-multiple
298 pass tests/userprog/args-many
299 pass tests/userprog/args-dbl-space
300 pass tests/userprog/sc-bad-sp
301 pass tests/userprog/sc-bad-arg
302 pass tests/userprog/sc-boundary
303 pass tests/userprog/sc-boundary-2
304 pass tests/userprog/halt
305 pass tests/userprog/exit
306 pass tests/userprog/create-normal
307 pass tests/userprog/create-empty
308 pass tests/userprog/create-null
309 pass tests/userprog/create-bad-ptr
310 pass tests/userprog/create-long
311 pass tests/userprog/create-exists
312 pass tests/userprog/create-bound
313 pass tests/userprog/open-normal
314 pass tests/userprog/open-missing
315 pass tests/userprog/open-boundary
316 pass tests/userprog/open-empty
317 pass tests/userprog/open-null
318 pass tests/userprog/open-bad-ptr
319 pass tests/userprog/open-twice
320 pass tests/userprog/close-normal
321 pass tests/userprog/close-twice
322 pass tests/userprog/close-stdin
323 pass tests/userprog/close-stdout
324 pass tests/userprog/close-bad-fd
325 pass tests/userprog/read-normal
326 pass tests/userprog/read-bad-ptr
327 pass tests/userprog/read-boundary
328 pass tests/userprog/read-zero
329 pass tests/userprog/read-stdout
330 pass tests/userprog/read-bad-fd
331 pass tests/userprog/write-normal
332 pass tests/userprog/write-bad-ptr
333 pass tests/userprog/write-boundary
334 pass tests/userprog/write-zero
335 pass tests/userprog/write-stdin
336 pass tests/userprog/write-bad-fd
337 pass tests/userprog/exec-once
338 pass tests/userprog/exec-arg
339 pass tests/userprog/exec-multiple
340 pass tests/userprog/exec-missing
341 pass tests/userprog/exec-bad-ptr
342 pass tests/userprog/wait-simple
343 pass tests/userprog/wait-twice
344 pass tests/userprog/wait-killed
```


345 pass tests/userprog/wait-bad-pid
346 pass tests/userprog/multi-recurse
347 pass tests/userprog/multi-child-fd
348 pass tests/userprog/rox-simple
349 pass tests/userprog/rox-child
350 pass tests/userprog/rox-multichild
351 pass tests/userprog/bad-read
352 pass tests/userprog/bad-write
353 pass tests/userprog/bad-read2
354 pass tests/userprog/bad-write2
355 pass tests/userprog/bad-jump
356 pass tests/userprog/bad-jump2
357 pass tests/userprog/no-vm/multi-oom
358 pass tests/filesys/base/lg-create
359 pass tests/filesys/base/lg-full
360 pass tests/filesys/base/lg-random
361 pass tests/filesys/base/lg-seq-block
362 pass tests/filesys/base/lg-seq-random
363 pass tests/filesys/base/sm-create
364 pass tests/filesys/base/sm-full
365 pass tests/filesys/base/sm-random
366 pass tests/filesys/base/sm-seq-block
367 pass tests/filesys/base/sm-seq-random
368 pass tests/filesys/base/syn-read
369 pass tests/filesys/base/syn-remove
370 pass tests/filesys/base/syn-write
371 All 76 tests passed.

372
373
374
375

376 Appendix

377

378 A Visit to PintOS Corporation

379 PintOS Corporation (PintCorp) employs a lot of employees and there is a significant turn-over of
380 these employees. New employees are contracted and old leave PintCorp regularly.

381 The company has a famous coffee house, where employees come often during their work to
382 rest, relax and of course to drink. Talking business here is an absolute no-no. The coffee house
383 has three doors. One of the door at the front is the entrance for the new employees. Each new
384 employee gets to enjoy a coffee break before work begins. There is a separate door to let the
385 employees who were away come back to work. They too get to enjoy their coffee before
386 resuming duties. These two doors are for entry only – no one is allowed out of these doors.
387 Employees at work enter and exit coffee house through the third door.

388 Employees spend as much time as they like in the coffee house. They can come in any time and
389 leave any time. An employee, who is not new to the company, goes to his work-desk and
390 resumes work diligently. The employees work till the work is complete or they want to have a
391 coffee or they need to go out of the Corporate area.

392 The arrangements for the new employee are practical. A new employee does not have a desk
393 to work from. They are given directions to the desk store. That is where they report to start
394 their work at PintCorp. The store gives the new employee a desk to work from. The new
395 employee sets their desk up and begin working.

396 PintCorp has no manager! Every employee follows the rules and work with the corporation till
397 the work is finished. The existing employees are able to hire new employees into PintCorp. A
398 final fact about PintCorp is that no more than one employee ever works at a time. Others are
399 either out of office or having a coffee break.

400 The employee who hires a new employee may wait for the hired employee to finish employee
401 to finish work before returning to work. But, some employees continue to work without waiting
402 for the hired employees to finish their assigned works.

403 If you have read the story carefully, you know that PintCorp calls its employees threads. Each
404 has a number `tid`. The coffee shop is called `THREAD_READY`. The working employee is
405 termed `THREAD_RUNNING`. And, those away are known as `THREAD_BLOCKED`.

406

407 Contributing Authors:

408 Vishv Malhotra, Gautam Barua, Rashmi Dutta Baruah